





UNIVERSITÀ
DEGLI STUDI DI BARI
ALDO MORO

DIPARTIMENTO DI
INFORMATICA



NLP project presentation

Sentiment analysis on movie reviews



Antonio Silletti

Elio Musacchio



Challenge description

Brief description of the problem



SVC approach

Simple technique using SVC classifier



Transformers approach

HuggingFace transformers with hyperparameters search



SBERT approach


Different classification techniques using SBERT embeddings



Neural network

Classification using custom Neural Network

Table of Contents






01

**Challenge
description**




Challenge description

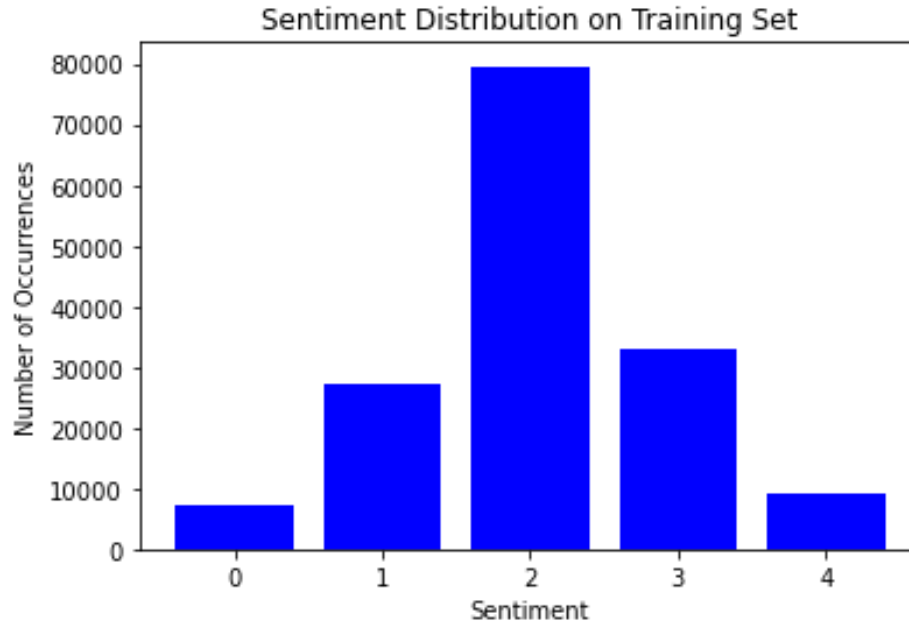



The competition consists in labeling phrases (from the Rotten Tomatoes dataset) based on the sentiment associated to them. Five values are possible for the sentiment labels:

- 0 -> negative
- 1 -> somewhat negative
- 2 -> neutral
- 3 -> somewhat positive
- 4 -> positive




Unbalanced Dataset



Dataset is unbalanced
towards the ***neutral***
sentiment

Inconsistency in labeling



- The fullstop at the end lowers the sentiment?



77	2	is worth seeking	.	<u>3</u>
78	2	is worth seeking		<u>4</u>

- Maybe not...

101	3	would have a hard time sitting through this one	.	<u>1</u>
102	3	would have a hard time sitting through this one		<u>0</u>

Sentiment Analysis is a difficult task even for humans, but it is even more difficult on this particular dataset!

02 SVC approach






Experimental protocol

HoldOut technique with 80% of phrases
«held» for the train set

*(naïve approach to obtain a first
feedback)*

SVC approach




- Preprocessing via Spacy:
 1. Word tokenization (with tokens also being lowercased in the process)
 2. Lemmatization
 3. NER


The first two operations are trivial (they are usually applied in order to reduce dimensionality of the vector space).

As for NER, the basic idea was to replace movie titles, actor's name, director's name, etc. with their respective tokens, as to avoid the occurrence of words such as "good" or "bad" that do not influence the Sentiment of the phrase

13216 569 The first question to ask about Bad Company



SVC approach



Once preprocessed, we used the CountVectorizer method of sklearn in order to turn each phrase into its sparse vector representation. Furthermore, we have set CountVectorizer to consider the occurrences of each word within a sentence only once.

Because in Sentiment Analysis task we are more interested in the occurrence of a word rather than its frequency in the corpus

Once obtained the vector representation of each phrase, we fed them to an SVC classifier with 'rbf' kernel, obtaining the following result on the test set:

YOUR RECENT SUBMISSION



svc_approach.csv

Submitted by UNIBA_NLP2122_Leshi · Submitted just now

Score: 0.62061





03

Transformers approach

With hyperparameters search via Ray Tune





Experimental protocol


Hyperparameters
search

→ HoldOut technique with 80% of phrases
«held» on 20% of the entire dataset

Training

→ Kfold with $K=5$ for training and validation

Transformers approach




We decided to try **transformers** because, differently from the previous approach, embeddings obtained by a hugging face model consider *contextual information*.


Due to the above, performing preprocessing operations it's not advised, aside from the tokenization operation (which is mandatory) performed by the tokenizer associated to the chosen model.

Thanks to the addition of context, we expected better results

(and we obtained them!)



Transformers approach




We choose the following models, based on results computed on 10% of the entire dataset:

- *bert-base-uncased*
- *gpt2*
- *albert-base-v2*


We also experimented with other models (such as *roberta-base*), but we obtained slightly worse results.

The models' choice was also guided by their size:

- We wanted to perform a comparison between a small model (*albert-base-v2*), a medium one (*bert-base-uncased*) and a large one (*gpt2*)



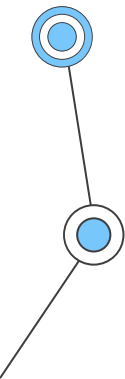
Transformers approach




Obviously, we needed to fine tune on a downstream-task each of the chosen models.

Since Hugging Face allows to attach any head (the module responsible for obtaining results) to the models it provides, we were able to attach the **SequenceClassification** head to models trained for a different task such as gpt2, which is used for text generation tasks.

We will now see the pipeline performed to fine tune each model.



Hyperparameters search




We searched for best hyperparameters using only 20% of the whole dataset (in order to reduce computational effort).

- The assumption is that the hyperparameters found are the best ones also for the rest of the dataset


The scheduler used is the **Population Based Training (PBT)**:

PBT takes its inspiration from genetic algorithms where each member of the population can exploit information from the remainder of the population

As the training of the population of neural networks progresses, parameters of the best models are **copied** to worse performing models, and at the same time they are **randomly perturbed**. This is done in order to hopefully obtain even **better results** than “previous generation” models.




PBT visualized



<https://docs.ray.io/en/latest/tune/tutorials/tune-advanced-tutorial.html>


Hyperparameters search




- Search space

```
tune_config = {  
    # search space  
    "per_device_train_batch_size": tune.choice([4, 8, 16, 32, 64]),  
    "num_train_epochs": tune.choice([2, 3, 4, 5]),  
    "seed": tune.randint(0, 43),  
    "weight_decay": tune.uniform(0.0, 0.3),  
    "learning_rate": tune.uniform(1e-4, 5e-5),  
    "lr_scheduler_type": tune.choice(['linear', 'cosine', 'polynomial', "cosine_with_restarts"]),  
}
```

- PBT also requires the definition of the Perturbation space (how the above parameters can be perturbed):
in our case it's an exact copy of the search space




Training



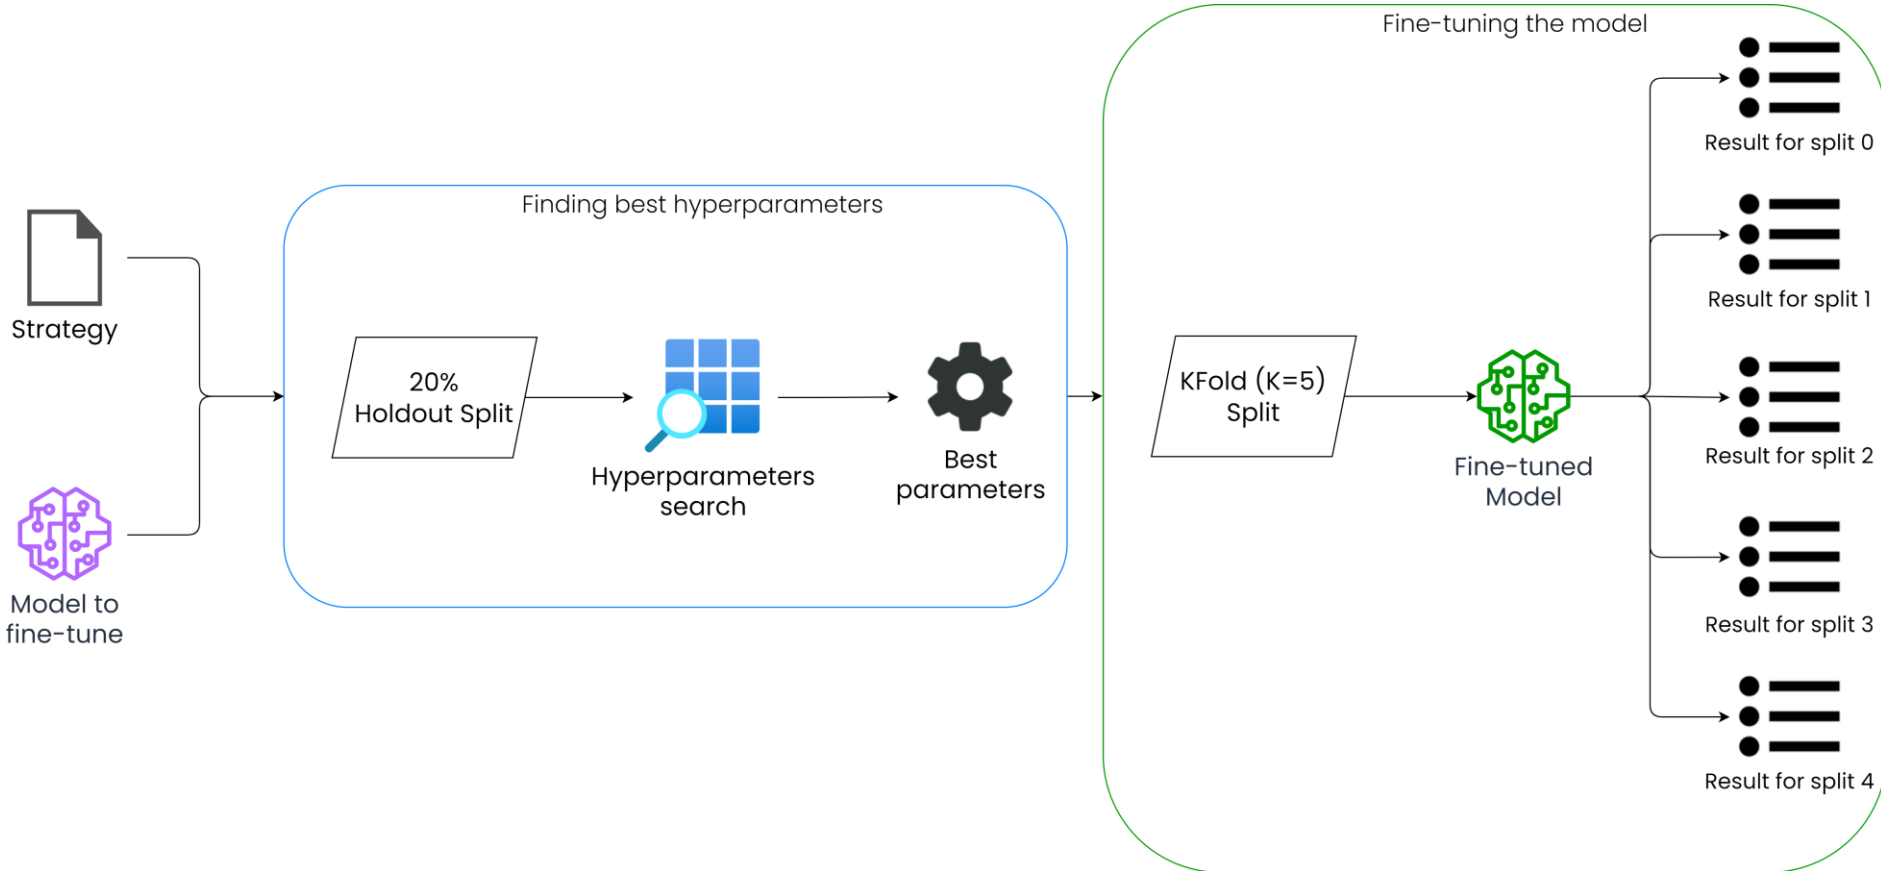
Each model is fine-tuned following each of the following strategies:

- **“Only phrase”** strategy: the single phrase is fed into the model
- **“With reference”** strategy: the single phrase and the original sentence from which the phrase comes from are fed into the model
- **“With POS”** strategy: the single phrase and its pos tags are fed into the model


We performed KFold partitioning with $K=5$ and we computed results for each split for each strategy.



Whole training visualized



Training



Due to the computational cost of the whole process, we rented a remote machine to perform the training:

A screenshot of a Vast.ai cloud instance details page. The interface shows various hardware and software specifications for a rented machine. At the top left, there's an ID '4632133', a machine name 'm:5359', and a host ID 'host:37035'. The Vast.ai logo is present. The main section highlights '1x RTX A4000' with '25.8 TFLOPS' and 'Max CUDA: 11.4'. Other specs include '16.1 GB' of memory at '447.5 GB/s', 'Z590 WIFI GUND...' motherboard with 'PCIe 4.0, 16x' slots at '22.4 GB/s', '11th Gen Core™ i...' CPU with '16.0/16 cpu' and '48/48 GB' cache, and 'nvme' storage at '2132 MB/s' with '60.8 GB' capacity. Network speeds are shown as '↑189.8 Mbps' and '↓199.1 Mbps'. The instance is marked as 'unverified' and has an 'Age: 1h 15m'. On the right, there are three buttons: 'STOP...', 'DESTROY...', and 'CONNECT'. At the bottom, it shows 'GPU: 96.0%, 61.0C' and 'Status: success, running pytorch/pytorch/jupyter'.

4632133 m:5359 host:37035 **unverified**

V **1x RTX A4000** **25.8** TFLOPS **16.1 GB**
vast.ai **Max CUDA: 11.4** **447.5 GB/s**

Z590 WIFI GUND... **↑189.8 Mbps**
PCIe 4.0, 16x **22.4 GB/s** **↓199.1 Mbps**

11th Gen Core™ i... **nvme**
16.0/16 cpu **48/48 GB** **2132 MB/s** **60.8 GB**


GPU: 96.0%, 61.0C Status: success, running pytorch/pytorch/jupyter

STOP...
DESTROY...
CONNECT


Even with this powerful machine, the complete training took roughly 2 days

Results BERT – 2 epochs


«Only Phrase» strategy



«With reference» strategy



«With POS» strategy




BEST results:


- «Only phrase» -> 0.7012
- «With reference» -> 0.7052
- «With POS» -> 0.7072

Results GPT2 – 2..3 epochs


«Only Phrase» strategy



«With reference» strategy



«With POS» strategy




BEST results:


- «Only phrase» -> 0.6975
- «With reference» -> **0.7003**
- «With POS» -> 0.6982

Results ALBERT – 2.3 epochs


«Only Phrase» strategy



«With reference» strategy




«With POS» strategy



BEST results:

- «Only phrase» -> **0.7001**
- «With reference» -> **0.51**
- «With POS» -> **0.6426**


TOP-5 Results on the test set



Submission and Description	Private Score	Public Score
bert-base-uncased_with_pos_split_2.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	<u>0.69572</u>	<u>0.69572</u>
bert-base-uncased_with_pos_split_0.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69469	0.69469
bert_with_pos_split3.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69300	0.69300
bert_with_pos_split1.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69237	0.69237
best_cm_1less_conv2d.csv 7 hours ago by UNIBA_NLP2122_Leshi add submission details	0.69224	0.69224




Final considerations





- Albert performed worse than the others (as expected) but trained much faster and produced smaller models w.r.t GPT2 and BERT
 - *1.6 gb for 5 splits vs 10 gb for 5 splits*
- GPT2 was the most consistent model across all the strategies, with its best result on the test set being **0.68975**
- The clear winner is BERT uncased with POS strategy: as a matter of fact, the models generated for each split with this combination performed better than all the other models on any strategy, and 4 out of 5 splits broke the 0.7 accuracy wall on the validation set

All graphs are visible at the following [link](#)



04

SBERT approach



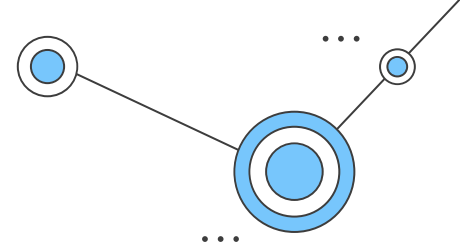


Experimental protocol

HoldOut technique with 80% of phrases
«held» for the train set

*(naïve approach to obtain a first
feedback)*


SBERT approach




We also performed some experiments on embeddings generated using the **sentence_transformers** library.

In this case, we used the pre-trained model *all-mpnet-base-v2* which, according to SBERT documentation, provides the best quality of results.


No preprocessing operations were applied since the model required the original textual data and no fine-tuning operation was performed this time.




TOP -5 SBERT PRETAINED MODELS



Model Name	Performance Sentence Embeddings (14 Datasets) ⓘ	Performance Semantic Search (6 Datasets) ⓘ	⚙ Avg. Performance ⓘ	Speed ⓘ	Model Size ⓘ
all-mpnet-base-v2 ⓘ	69.57	57.02	63.30	2800	420 MB
multi-qa-mpnet-base-dot-v1 ⓘ	66.76	57.60	62.18	2800	420 MB
all-distilroberta-v1 ⓘ	68.73	50.94	59.84	4000	290 MB
all-MiniLM-L12-v2 ⓘ	68.70	50.82	59.76	7500	120 MB
multi-qa-distilbert-cos-v1 ⓘ	65.98	52.83	59.41	4000	250 MB



How we used SBERT embeddings




Three different techniques were considered using SBERT-generated embeddings:

1. Clusters represented by centroids of training embeddings with the same label and classification via cosine similarity (picking the label of the cluster's centroid with highest similarity);
2. Classification via Multi Layer Perceptron;
3. Classification via Dense Neural Network.

The first two techniques didn't provide satisfactory results during the experimental protocol after different trials (the obtained accuracy on the validation set was less than 0.6).

Using a dense neural network for classification the accuracy reached 0.66 top on the test set.






05

Custom
NN




Why Custom NN?




We inspected the HEAD that performs text classification provided by HuggingFace, and it basically only performs a **dropout** followed by a **dense layer** on the pooled output of the 13 layers of the encoders.

We tried to attach a much more complex Custom Head that performs classification

We tested our Custom Head on embeddings generated by SBERT (by using the all-mpnet-base-v2 model) and bert-base-uncased with POS strategy (the model which gave us the best results on the Transformers approach)



Custom NN architecture




Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 13, 1, 384]	1,534
LeakyReLU-2	[-1, 13, 1, 384]	0
Conv2d-3	[-1, 13, 1, 192]	1,534
LeakyReLU-4	[-1, 13, 1, 192]	0
BatchNorm2d-5	[-1, 13, 1, 192]	26
Conv2d-6	[-1, 13, 1, 96]	1,534
LeakyReLU-7	[-1, 13, 1, 96]	0
BatchNorm2d-8	[-1, 13, 1, 96]	26
Conv2d-9	[-1, 13, 1, 48]	1,534
LeakyReLU-10	[-1, 13, 1, 48]	0
BatchNorm2d-11	[-1, 13, 1, 48]	26
Conv2d-12	[-1, 13, 1, 24]	1,534
LeakyReLU-13	[-1, 13, 1, 24]	0
BatchNorm2d-14	[-1, 13, 1, 24]	26
Dropout-15	[-1, 13, 1, 24]	0
Flatten-16	[-1, 312]	0
Linear-17	[-1, 5]	1,565

- Input is the embedding obtained by bert-uncased averaged to get a single embedding for the sentence ($dim=2$)
- The idea is to use all the layers of the encoders ($dim=1$) and not pool them or using only the last layer (*information bottleneck*)






Custom NN best-result



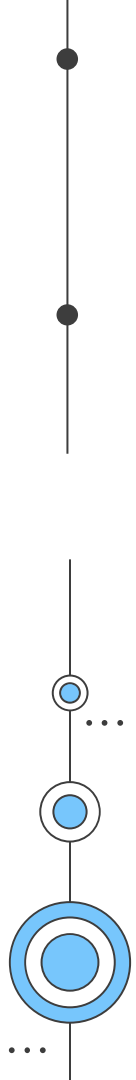
Submission and Description	Private Score	Public Score
bert-base-uncased_with_pos_split_2.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69572	0.69572
bert-base-uncased_with_pos_split_0.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69469	0.69469
bert_with_pos_split3.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69300	0.69300
bert_with_pos_split1.csv 2 days ago by UNIBA_NLP2122_Leshi add submission details	0.69237	0.69237
best_cm_1less_conv2d.csv 7 hours ago by UNIBA_NLP2122_Leshi add submission details	0.69224	0.69224

5-th best
result
(with bert
embeddings)

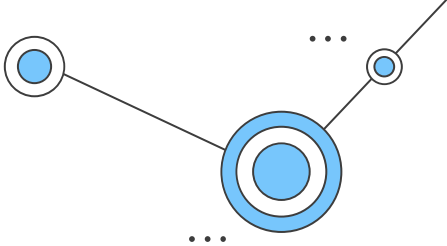













Leaderboard results



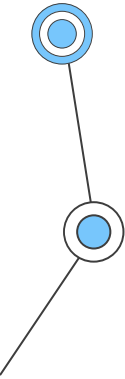
Old leaderboard




#	Team	Members	Score
1	Mark Archer		0.76526
2	Armineh Nourbakhsh		0.76096
3	Merlion	   	0.70936
4	Puneet Singh		0.70789
5	Yoon		0.68765
6	DrStrangelove		0.67931

Our result:
0.69572

With our best result we would be 5-th over 861 total groups
in the old leaderboard



Public leaderboard (more recent)



All Your Work Shared With You Bookmarks

Best Score ▾



Fastai with 🤖 Transformers (BERT, RoBERTa, ...)

Updated 2Y ago

Score: 0.70071 · 79 comments · Sentiment Analysis on Movie Reviews

▲ 307

🥇 Gold ...



bert_experiment

Notebook copied with edits from Nikita Sharma · Updated 3Y ago

Score: 0.6978 · 0 comments · Sentiment Analysis on Movie Reviews

▲ 3

...



BERT for Sentiment Analysis - 5th Place Solution

Updated 10mo ago

Score: 0.69644 · 2 comments · Sentiment Analysis on Movie Reviews +1

▲ 13

🥉 Bronze ...



Sentiment Analysis On Movie Reviews

Updated 6mo ago

Score: 0.69572 · 0 comments · Sentiment Analysis on Movie Reviews

▲ 0

...

*Tied with
our result:*
0.69572

With our best result we would be tied for 4-th place over ?
total groups in the public and more recent leaderboard



A decorative network diagram consisting of several blue circular nodes connected by thin grey lines. The nodes are arranged in a non-linear fashion, with some having concentric circles. Ellipses (...) are placed near some nodes to indicate a larger network. The diagram is positioned around a central light blue cloud-like shape.

Thanks!